



a

Matched Q3 segments

**b**

Scale trend (Q3 SPJ)



Q3 SPJ, $p=0.5$. FAIR uncontended 3-run: ProvSQL sieve (its practical exact method) vs ours end-to-end (construct + shared ROBDD compile + WMC, $O(N)$ ordering). Parity EXACT (both $0.5^3=0.125$). Ours 2.6–2.9x faster across SF 0.01/0.1/0.3 — stock engine, no fork. The gap is in-memory batch WMC vs ProvSQL's per-answer in-database evaluation; method-sensitive (its d4 isn't batched here).