



a**b**

Q3 SPJ PQE, per-token $p=0.5$. Probability parity is EXACT (both compute $0.5^3=0.125$ per answer). ProVSQL = honest per-answer probability_evaluate (materialized, not the count(*) pruning artifact), ~ 0.24 ms/answer; ours (stock engine + client compiler, no engine fork) is faster at every scale (2.4 $\times$ at SF 0.1, 1.5 $\times$ at SF 0.3).