



a

Reconvergent (ours faster, both)

**b**

Scale trend (Q3 SPJ)



p=0.5, FAIR uncontended (≥ 3 -run median). ProvsSQL sieve (its practical exact method) vs ours end-to-end (construct + shared ROBDD compile + WMC, $O(N)$ ordering). Parity EXACT. Ours faster at every scale: reconvergent 2.3–2.5 $\times$ (a), Q3 SPJ 2.6–2.9 $\times$ across SF 0.01/0.1/0.3 (b) — stock engine, no fork. The gap is in-memory batch WMC vs ProvsSQL's per-answer in-database evaluation; method-sensitive (its d4 isn't batched here).